

ST502 Project

Halid Kopanski and Justin Feathers

Part I

Framingham Heart Study

The Framingham data set contains the diastolic blood pressure of 300 smokers and nonsmokers. For this study we will assume the following:

$$H_0: \mu_{ns} - \mu_s = 0$$

$$H_A: \mu_{ns} - \mu_s \neq 0$$

The null hypothesis being there is no difference in the blood pressure between smokers and non smokers. We believe that there is a difference and this paper will prove if we have enough evidence to reject the null hypothesis.

Normality of Data

Plotting the density of total data (see Appendix I: Figure 1A), we can that it closely follows a normal distribution, Density plots of the split data exhibit similar forms (see Appendix I: Figure 1B). We can also calculate the kurtosis and the skew of the data, which are 0.812, 0.88 for kurtosis and skew respectively, and see that they are quite close to zero. Therefore, we will assume the data and the subsequent split data to be normal.

Statistical Analysis

In the first analysis we will be assuming equal variances, thereby allowing us to use pooled sample variance (see Appendix I: Equation 1).

The pooled sample variance of the data is 510 using a degree of freedom value 298. In this case, we reject the null hypothesis because the calculated p value, 0.0041, is smaller than the chosen α value of 0.05. In terms of t values, our observed t value of -3.04 is smaller than the t value, -1.97 for a two sided α of 0.05.

When computing the observed t value using the assumption that the population variances are not equivalent (variance smoker is 352.2 and variance nonsmoker is 562.1), a value of -2.9 is obtained. Comparing that the two sided α of -1.96 with 693 degrees of freedom (as computed using the Satterthwaite Approximation)....

The observed 95% confidence intervals are -15.08 to -3.23 for pooled sample variance and -15.36 to -2.95 for non pooled sample variance. It can be seen that 0 does not fall into the 95% confidence interval in either case.

In all three case, there is sufficient evidence to reject the null hypothesis and to support the alternative at an α level of 0.05.

Part II

The number of scenarios where the null was rejected is 890.

Appendix I: Equations and Figures

Figures




Figure 1: Data plots

Equations:

Equation 1: Pooled Sample Variance:

$$S_p^2 = \frac{(n_{ns}-1)S_{ns}^2 + (n_s-1)S_s^2}{n_{ns} + n_s - 2}$$

Equation 2: Observed T Statistic (pooled variance):

$$t_{obs} = \frac{\mu_{ns} - \mu_s}{S_p \sqrt{\frac{1}{n_{ns}} + \frac{1}{n_s}}}$$

Equation 3: Observed T Statistic (distinct variance):

$$t_{obs} = \frac{\mu_{ns} - \mu_s}{\sqrt{\frac{S_{ns}^2}{n_{ns}} + \frac{S_s^2}{n_s}}}$$

Equation 4: Satterthwaite Approximation:

$$\nu = \frac{(\frac{S_{ns}^2}{n_{ns}} + \frac{S_s^2}{n_s})^2}{\frac{(S_{ns}^2/n_{ns})^2}{n_{ns}-1} + \frac{(S_s^2/n_s)^2}{n_s-1}}$$

Appendix II: Results

Data Mean: 134.935

Data Sample Variance: 524.0852258

Smoker Mean: 128.0666667

Smoker Variance: 352.2117117

Nonsmoker Mean: 137.2244444

Nonsmoker Variance: 562.1447123

Difference in mean: -9.1577778

Pooled Sample Variance: 510.0136987

Nonpooled Sample Variance: 9.9936838

CI Pooled: -15.08, -3.23

CI Nonpooled: -15.36, -2.95

Appendix III: Code

```
knitr::opts_chunk$set(echo = TRUE)

#Use required packages
library(tidyverse) #for plots and data manipulation
library(cowplot) #aligning plots
library(e1071) #skew and kurtosis calculations only

df_data <- read_csv("framingham_data.csv") # Read in data
df_data$index <- seq(nrow(df_data)) # Add an index column

#df_data %>% summary # Summarize Data

# Split data into smoker and nonsmoker
df_smoker <- df_data %>% filter(currentSmoker == 1)
df_nonsmoker <- df_data %>% filter(currentSmoker == 0)

#Plot and compare split data

#options(repr.plot.width = 6, repr.plot.height = 4, repr.plot.res = 150)

plot_colors <- c("#001427", "#708d81", "#f4d58d", "#bf0603", "#8d0801")
y_limits <- c(0, 0.0225)

total_data <- ggplot(df_data) + geom_density(aes(sysBP),
                                             fill = plot_colors[1],
                                             alpha = 0.6) +
  ylim(y_limits) + ylab("Density") + xlab("")

sep_data <- ggplot() + geom_density(data = df_smoker, aes(sysBP),
                                   fill = plot_colors[3], alpha = 0.6) +
  geom_density(data = df_nonsmoker, aes(sysBP),
               fill = plot_colors[5], alpha = 0.6) +
  ylim(y_limits) + ylab("") + xlab("Systolic Blood Pressure")

plot_3 <- ggplot() + geom_density(data = df_smoker, aes(sysBP),
                                  fill = plot_colors[3], alpha = 0.5) +
  geom_density(data = df_data, aes(sysBP),
               fill = plot_colors[1], alpha = 0.5) +
  geom_density(data = df_nonsmoker, aes(sysBP),
               fill = plot_colors[5], alpha = 0.5) +
  ylim(y_limits) + ylab("") + xlab("")

#plot_grid(total_data, sep_data, plot_3, align = 'vh',
#           hjust = -1, nrow = 1, ncol = 3)

data_kurtosis <- kurtosis(df_data$sysBP)
data_skew <- skewness(df_data$sysBP)

#Create a sample variance function to ensure proper calculation
sample_variance <- function(x, sample = TRUE){
  if (sample == TRUE){
    sum((x - mean(x))^2) / (length(x) - 1)
  } else if (sample == FALSE) {
```

```

    sum((x - mean(x))^2) / (length(x))
  }
}
#Create pooled sample variance function
f_pooled_variance <- function(x, y){
  ((length(x) - 1) * sample_variance(x) +
   (length(y) - 1) * sample_variance(y)) /
  (length(x) + length(y) - 2)
}

# Create a Satterthwaite Approximation Function

satterth <- function(s1, s2, n1, n2){
  term1 <- s1/n1
  term2 <- s2/n2
  nu <- (term1 + term2)^2 / ((term1/(n1 - 1)) + (term2/(n2 - 1)))
  return(floor(nu))
}

# Common values for analysis

alpha <- 0.05

mu_smoker <- mean(df_smoker$sysBP)
var_smoker <- sample_variance(df_smoker$sysBP)
n_smoker <- length(df_smoker$sysBP)

mu_nonsmoker <- mean(df_nonsmoker$sysBP)
var_nonsmoker <- sample_variance(df_nonsmoker$sysBP)
n_nonsmoker <- length(df_nonsmoker$sysBP)

# Two Sample T-test - Pooled Sample Variance - P-value

dof_1 <- (n_smoker + n_nonsmoker - 2)

p_sample_var_1 <- f_pooled_variance(df_smoker$sysBP,
                                   df_nonsmoker$sysBP)

t_obs_1 <- (mu_smoker - mu_nonsmoker) / (sqrt(p_sample_var_1) * sqrt(1/n_smoker + 1/n_nonsmoker))

t_stat_1 <- qt(alpha / 2, dof_1)

p_value_obs_1 <- dt(t_obs_1, dof_1)

#Two Sample T-test - Difference Variance Sample Variance - P-value

dof_2 <- satterth(var_smoker, var_nonsmoker, n_smoker, n_nonsmoker)

np_sample_var_2 <- (var_nonsmoker/n_smoker + var_nonsmoker/n_nonsmoker)

t_obs_2 <- (mu_smoker - mu_nonsmoker) / (sqrt(var_nonsmoker/n_smoker + var_nonsmoker/n_nonsmoker))

t_stat_2 <- qt(alpha / 2, dof_2)

p_value_obs_2 <- dt(t_obs_2, dof_2)

# Confidence Limits

```

```

diff_mu <- mu_smoker - mu_nonsmoker

#Pooled Sample variance

CL_pooled <- t_stat_1 * (sqrt(p_sample_var_1/n_smoker + p_sample_var_1/n_nonsmoker))

#Non pooled Sample variance

CL_nonpooled <- t_stat_2 * (sqrt(var_nonsmoker/n_smoker + var_nonsmoker/n_nonsmoker))

CI_pooled <- round(c(diff_mu + CL_pooled, diff_mu - CL_pooled), 2)

CI_nonpooled <-round(c(diff_mu + CL_nonpooled, diff_mu - CL_nonpooled), 2)

#Part II
set.seed(1)

alpha <- 0.05

#Create a sample variance function to ensure proper calculation
sample_variance <- function(x, sample = TRUE){
  if (sample == TRUE){
    sum((x - mean(x))^2) / (length(x) - 1)
  } else if (sample == FALSE) {
    sum((x - mean(x))^2) / (length(x))
  }
}

test_function <- function (x, y){

  #Bootstrap datasets
  sample1 <- sample(x, size = length(x), replace = TRUE)
  sample2 <- sample(y, size = length(y), replace = TRUE)

  dof <- length(x) + length(y) - 2 #Determine degrees of freedom

  #calculate the mean and variance of the two bootstrap sets
  mu_1 <- mean(sample1)
  var_1 <- sample_variance(sample1, sample = TRUE)

  mu_2 <- mean(sample2)
  var_2 <- sample_variance(sample2, sample = TRUE)

  #Calculate the pooled sample variance
  pooled_sample <- ((length(x) - 1) * var_1 +
                    (length(y) - 1) * var_2) / (length(x) + length(y) - 2)

  #calculate the observed t statistic
  ttest <- (mu_1 - mu_2) / (sqrt(pooled_sample/length(x) +
                                pooled_sample/length(y)))

  #Determine whether or not the null hypothesis can be
  #rejected (1 = rejected, 0 = not rejected)
  verdict <- !between(ttest, qt(alpha / 2, dof), qt(1 - alpha / 2, dof))

  #Return calculated values
  return(c(mu_1, var_1, mu_2, var_2, ttest, verdict))
}

```



```

}

set.seed(1980)

sim_data_results <- rep(0, 6)

for (i in seq(1, 1000)){
  #Run the boot strap a 1000 times and store the results in a matrix
  sim_data_results <- rbind(sim_data_results,
                           test_function(df_nonsmoker$sysBP,
                                         df_smoker$sysBP))
}

#Move data into dataframe for easier processing

df_sim_data <- data.frame(sim_data_results[2:nrow(sim_data_results),])

colnames(df_sim_data) = c("Mean Smoker", "Variance Smoker",
                          "Mean NonSmoker", "Variance NonSmoker",
                          "T statistic", "Null Reject")

#Calculate the number of scenarios where the null was rejected
no_reject_null <- sum(df_sim_data$'Null Reject')
# Plotting
options(repr.plot.width = 6, repr.plot.height = 4, repr.plot.res = 150)

plot_grid(total_data, sep_data, plot_3, align = 'vh',
          hjust = -1, nrow = 1, ncol = 3, labels = c("A", "B", "C"))

```